


Every closed iteration is a model version. Your run closed 12 iterations in 6 waves (K=2 per wave) before halting on a score plateau. Each one left a complete, versioned model on disk — its code, its config, its scores. One model, three names: an iteration number, a coordinate hash, a fingerprint.

THE SCOREBOARD — ALL 12 VERSIONS, REAL SCORES




ONE MODEL, THREE NAMES
exp_003 ← folder + iteration number
hash dee2d34d02e4 ← the coordinate (knob settings)
fpr b605fd4bc030... ← the built model itself
Coordinates can repeat ideas; fingerprints can't. Dedupe runs on fingerprints, so two versions are never secretly the same model.

WHERE VERSIONS LIVE

```
workspace/  
  experiments/exp_000 ... exp_011/ ← one folder per version  
    model.py · config.json · judge.json · sim_stdout.log  
  best_model.py ← copy of the selected champion (#3)  
  best_model_confirmed.py ← copy of the hidden-exam winner (#10)  
  deploy/ ← what ships: weights + runtime + ID card
```

THE LIFE OF ONE VERSION



THE FOUR GATES

- 1

ANCHORED SCORING

The Judge reads the sim log and scores 0–100, but must stay within **±3 of the deterministic sim_score** — otherwise the answer is rejected and retried. Opinions can't drift from the machine.
- 2

RUNNING BEST

After each iteration, Python compares: higher judge score wins; tie → higher sim_score. Iteration 3 (82 / sim 82.28) beat 5 (82 / 81.95) and 6 (82 / 81.89) on the tie-break. Every win copies the model to `best_model.py`.
- 3

TIE SET AT HALT

Picking the max of 12 noisy scores rewards luck (winner's curse). So at halt, every version within **2 standard errors** of the champion's sim_score becomes a candidate — here, 6 versions tied statistically.
- 4

THE HIDDEN EXAM

Each candidate is retrained and replayed on **200 fresh scenarios (seeds 95000+)** never used during selection. New data, same deterministic simulator. The result is written to `confirmation.json`.
- WHY TWO CROWNS?** The spec's answer to "which is best" must be the judge-score rule (reproducible by definition). The evidence answer comes from the hidden exam. The factory reports both — and tells you when they disagree.

CONFIRMATION.JSON — THE REAL RESULTS

version	judge	eval sim	hidden-exam sim	hidden return	Δ vs #3
exp_003	82	82.28	81.16	32.88	—
exp_005	82	81.95	80.99	32.84	−0.04
exp_006	82	81.89	80.40	32.70	−0.18
exp_010	82	81.50	81.37	32.93	+0.05
exp_002	81	81.39	80.91	32.82	−0.06
exp_011	79	81.16	81.21	32.89	+0.01


- THE TWIST**

The judge crowned **#3**. The hidden exam crowned **#10** (81.37 vs 81.16). The margin is small — 0.21 sim points, within noise — but that is exactly the point: the exam exists so the final answer is *measured*, not inherited from selection luck.
- WHICH FILE DO I USE?**

`best_model.py` — #3, the spec's answer (highest judge score, tie → sim).
`best_model_confirmed.py` — #10, the hidden-exam answer.
`deploy/` — ships **#10**: the evidence winner.
- READ THE NOTE**

"The selected best (highest judge score) is Iteration 003; the confirmed best on the hidden set is Iteration 010."

MEAN RETURN vs THE FIXED LANDMARKS (computed before any model trained)



CHAMPION CARD — exp_003

knob	setting
features_recipe	reference_35 (35 features)
hidden_layers	[64, 64] — 6,854 params
teacher	v2 (two-step lookahead)
target_recipe	soft, temperature 0.5
data_recipe	behaviour + self, 84 DAgger episodes
decision_calibration	noop_margin_down (−0.75 NOOP bias)
ensemble_seeds	3
optim_recipe	reference (adamw, lr 5e-4, 30 epochs)

success rate 98% overflow 0.0 val_acc 59.6% train≈val → no overfit

WHAT SHIPS — workspace/deploy/

deploy/

- intervention_policy.pt ← trained weights (iteration 10)
- policy_model.py ← the model class
- policy_runtime.py ← inference wrapper
- policy_card.json ← the ID card

THE ID CARD SAYS

iteration 10 · eval sim 81.5 · hidden-exam sim 81.37 · 35 named features · 6 actions · **parity**

check: all 50 replay turns identical between simulator and runtime · trained in the deterministic synthetic simulator (sim-to-real caveats in deploy/README.md).

THE RUN IN NUMBERS

12 versions · 6 waves · 0 poisoned · 57 LLM calls · **\$7.42 total** · halted by the plateau rule, not by a human.